

Statistical Methods for High Dimensional Biology

STAT/BIOF/GSAT 540

Lecture 10 – Linear Models Part IV

Amrit Singh

Feb 06 2016

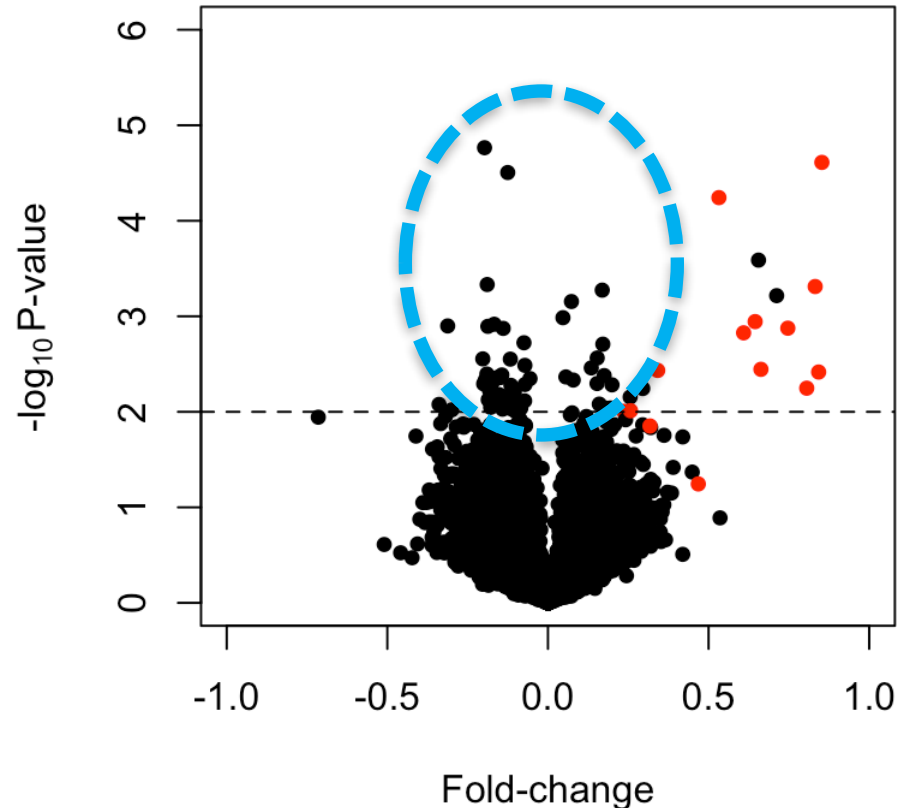
****Based on slides by Sara Mostafavi & Dr. Jennifer Bryan****

Motivation – data(“spikein95”)

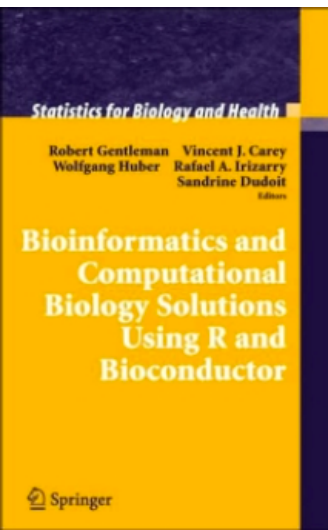
spikein95

- 2 groups of 3 samples each
- groups differ w/r to concentration of 16 probesets

A) Volcano plot for t-test



1470 differentially expressed genes!! –
majority have very small variances



outline

- Review
 - Linear regression framework
- Large scale differential expression analysis:
 - Assessing ALL genes (in a univariate way)
 - i.e., same model, except run it >20K times
 - Empirical Bayes → moderated test statistic
 - running Limma in R

$$Y = X\alpha + \varepsilon$$

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \begin{bmatrix} \alpha_0 \\ \alpha_1 \end{bmatrix} + \begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{bmatrix} = \begin{bmatrix} \alpha_0 \cdot 1 + \alpha_1 \cdot x_1 \\ \alpha_0 \cdot 1 + \alpha_1 \cdot x_2 \\ \vdots \\ \alpha_0 \cdot 1 + \alpha_1 \cdot x_n \end{bmatrix} + \begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{bmatrix} = \begin{bmatrix} \alpha_0 + \alpha_1 x_1 + \varepsilon_1 \\ \alpha_0 + \alpha_1 x_2 + \varepsilon_2 \\ \vdots \\ \alpha_0 + \alpha_1 x_n + \varepsilon_n \end{bmatrix}$$

$$y_i = \alpha_0 + \alpha_1 x_i + \varepsilon_i$$

Here we are just fitting a line but using matrix notation to handle all n observations at once, more elegantly.

Big pay-offs ensue

Industrial scale model fitting is good because things like this are not recomputed 30K times unnecessarily*

$Y = X\alpha + \varepsilon$ regression model

$\hat{\alpha} = (X^T X)^{-1} X^T Y$ the MLE and OLS estimator of α

$\hat{\sigma}^2 = \frac{1}{n-p} \hat{\varepsilon}^T \hat{\varepsilon}$ the estimated error variance

$\hat{V}(\hat{\alpha}) = \hat{\sigma}^2 (X^T X)^{-1}$ the estimated covariance matrix of $\hat{\alpha}$

How test $H_0 : \alpha_j = 0$?

With a t-statistic. Under H_0 , we have (at least approximately) that:

$$\frac{\hat{\alpha}_j}{\widehat{se}(\hat{\alpha}_j)} \sim t_{n-p}$$

so a p-value is obtained by computing a tail probability for the observed value of $\hat{\alpha}_j$ from a t_{n-p} distribution.

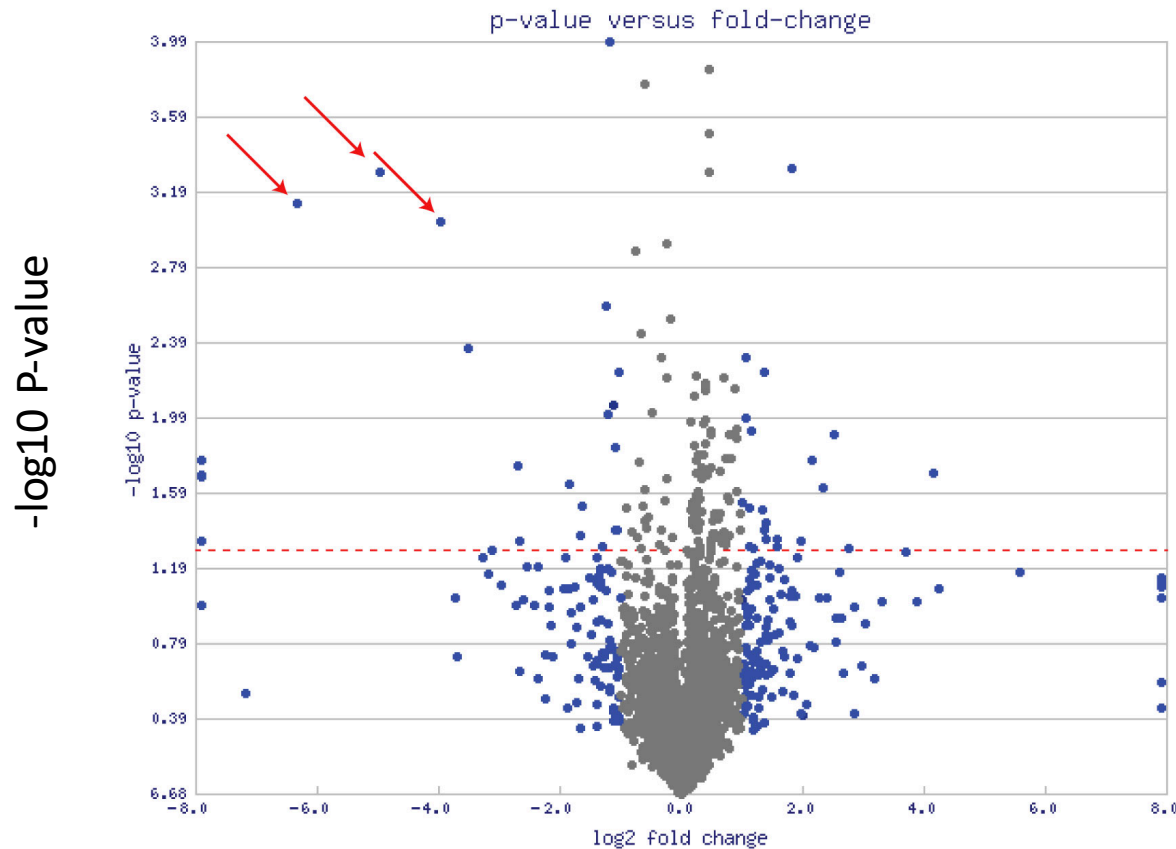
* under the hood, `lm()` is doing something more clever and numerically stable than this

Recurring theme in analysis of “high dimensional” biological data

Genes with very small p-values BUT subtle effects (small effect sizes)

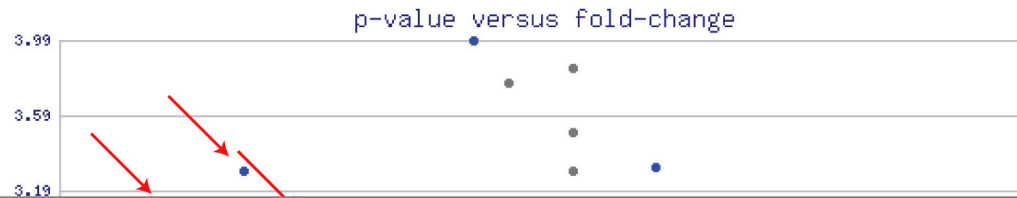
- Replication rate typically lower for genes with subtle effects
- Ad hoc filters: require small pvalues and large effect sizes

Observed (i.e., empirical) issues with the “standard” (i.e., t-test) approach for assessing differential expression

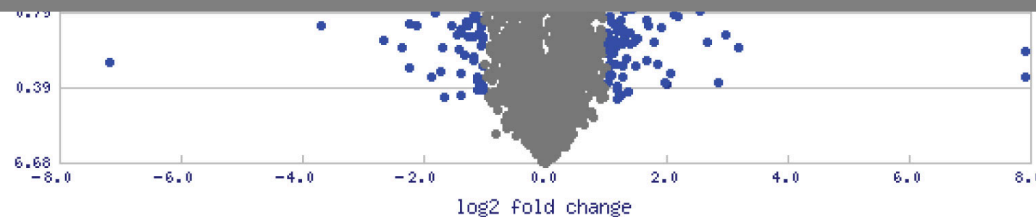


Log2 fold change (i.e., effect size)

Observed (i.e., empirical) issues with the “standard” (i.e., t-test) approach for assessing differential expression

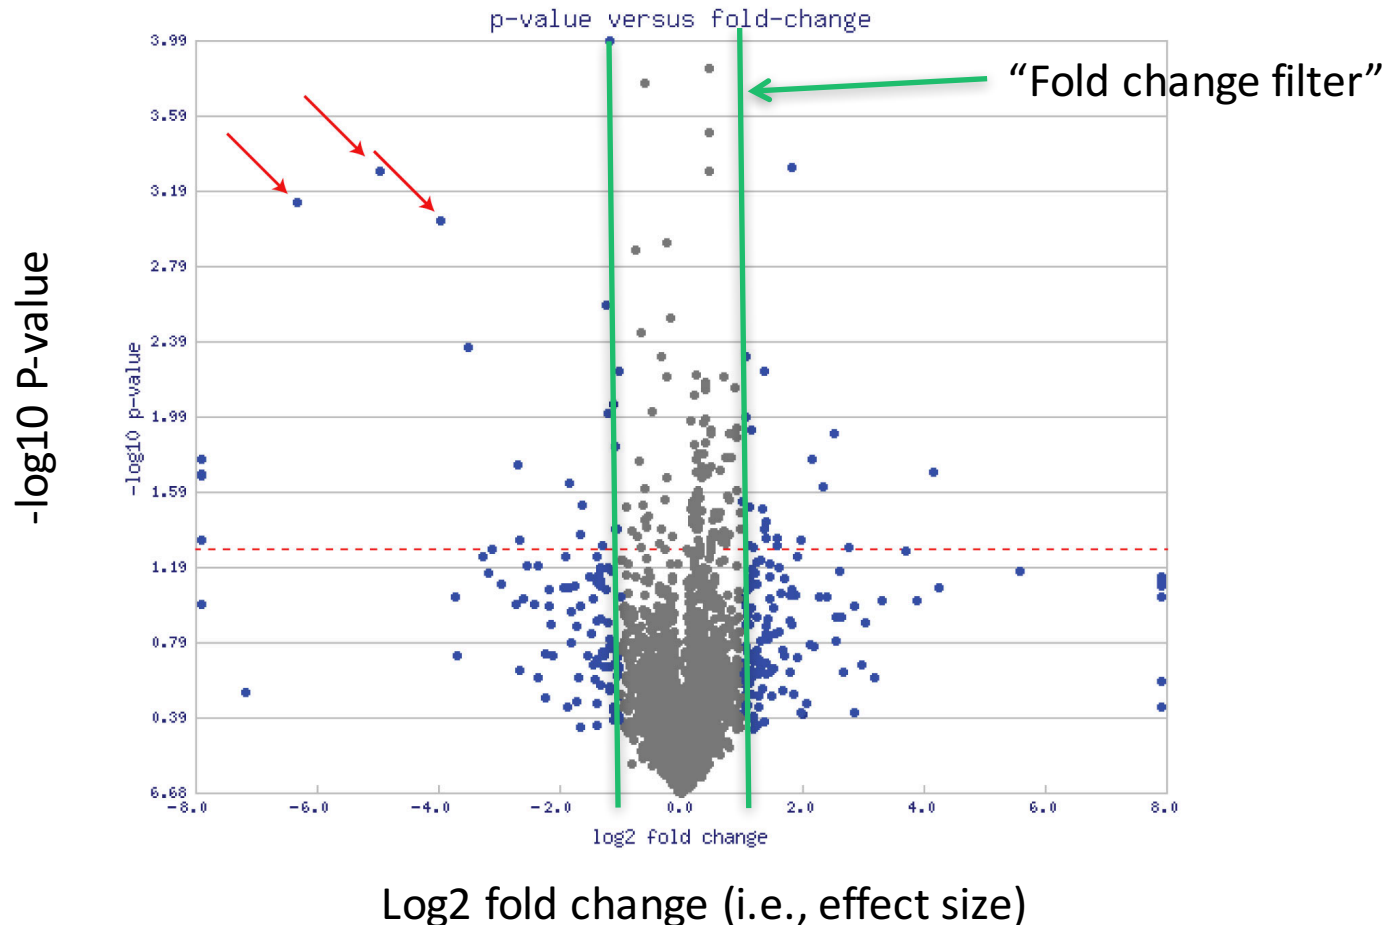


Some genes with very small pvalues (large $-\log_{10}$ pvalues) are not biologically meaningful.



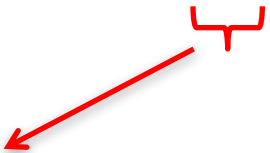
Log2 fold change (i.e., effect size)

Observed (i.e., empirical) issues with the “standard”
(i.e., t-test) approach for assessing differential expression

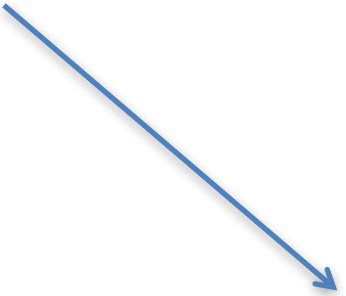


How do we end up with small pvalues but subtle effects?

$$t_{gj} = \frac{\hat{\alpha}_{gj}}{SE(\hat{\alpha}_{gj})} = \frac{\hat{\alpha}_{gj}}{s_g \sqrt{v_j}} \sim t_d \text{ under } H_0$$



Small variance leads
to large t stat,
leading to small p



d=residual degree of
freedom

Let's review how we derive the test statistics from our linear model

$$Y_g = X_g \alpha_g + \varepsilon_g$$

the “g” in the subscript reminds us that we'll be fitting a model like this for each gene g

most of the time the design matrices X_g are, in fact, the same for all g; I'm going to just use X

also, let's record the residual degrees of freedom:

$$d_g = d = n - \text{dimension of } \alpha$$

$$Y_g = X\alpha_g + \varepsilon_g \quad \text{the data model}$$

$$\text{var}(\varepsilon_g) = \sigma_g^2$$

$$s_g^2 = \frac{1}{n-p} \hat{\varepsilon}^T \hat{\varepsilon} \quad \text{estimated error variance (} p \text{ is the dimension of } \alpha \text{)}$$

$$\text{var}(\hat{\alpha}_g) = (X^T X)^{-1} s_g^2 = V s_g^2$$

“unscaled covariance”

$$Y_g = X\alpha_g + \varepsilon_g$$

$$\begin{bmatrix} y_{g1} \\ y_{g2} \\ \vdots \\ y_{gn_g} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ \vdots & \vdots & \vdots \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ \vdots & \vdots & \vdots \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ \vdots & \vdots & \vdots \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \mu_{g1} \\ \mu_{g2} \\ \mu_{g3} \end{bmatrix} + \begin{bmatrix} \varepsilon_{g1} \\ \varepsilon_{g2} \\ \vdots \\ \varepsilon_{gn_g} \end{bmatrix}$$

What would the estimated covariance matrix $\hat{V}(\hat{\alpha}_g) = s_g^2(X^T X)^{-1} = Vs_g^2$ actually look like in a concrete example?

$$\hat{V}(\hat{\alpha}_g) = \begin{bmatrix} \hat{V}(\hat{\mu}_1) & \widehat{\text{cov}}(\hat{\mu}_1, \hat{\mu}_2) & \widehat{\text{cov}}(\hat{\mu}_1, \hat{\mu}_3) \\ \widehat{\text{cov}}(\hat{\mu}_1, \hat{\mu}_2) & \hat{V}(\hat{\mu}_2) & \widehat{\text{cov}}(\hat{\mu}_2, \hat{\mu}_3) \\ \widehat{\text{cov}}(\hat{\mu}_2, \hat{\mu}_3) & \widehat{\text{cov}}(\hat{\mu}_2, \hat{\mu}_3) & \hat{V}(\hat{\mu}_3) \end{bmatrix}$$

So far, nothing new: the “regular” t statistics for gene g and parameters j :

$$t_{gj} = \frac{\widehat{\beta}_{gj}}{s_g \sqrt{v_j}} \sim t_d \text{ under } H_0$$

How do we end up with small p-values but subtle effects?

$$t_{gj} = \frac{\hat{\alpha}_{gj}}{SE(\hat{\alpha}_{gj})} = \frac{\hat{\alpha}_{gj}}{s_g \sqrt{v_j}} \sim t_d \quad \text{under } H_0$$

How would you modify the formula for t if you wanted to avoid having small p's for genes with subtle effects?

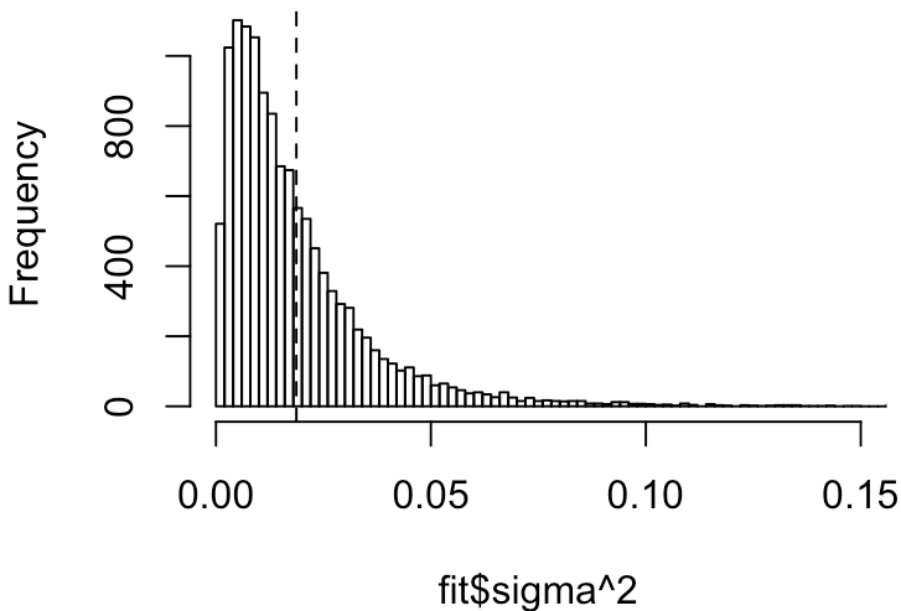
We could even come up with “weights w ” so to take a weighted average between estimated var (s_g) and *prior* var (s_0).

$$t_{gj} = \frac{\hat{\alpha}_{gj}}{(w_0 s_0 + w_1 s_g) \sqrt{v_j}}$$

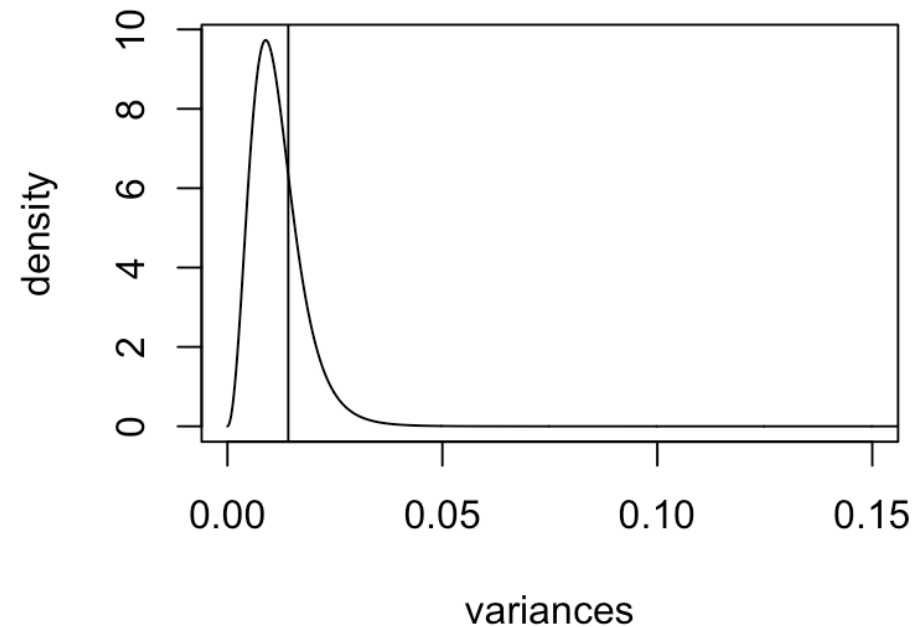
Moderated t statistics → Derived using the Empirical Bayes framework.

Modelling the distribution of all genewise variances

`mean(fit$sigma^2)=0.019`



scaled F-distribution,
`s_0=0.014, d_0=7.6, d_g=4`



Fit scaled F-distribution

$$s^2 \sim s_0^2 F_{d_g, d_0}$$



Bayes Theorem

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

Bayes Theorem for a continuous random variable

Posterior distribution of
parameters given data

Likelihood function of
data given a set of
parameters

$$P(\theta | x) = \frac{P(x | \theta)P(\theta)}{P(x)}$$

Prior distribution of
parameters

where

Marginal distribution
of data over all
parameter values

$$P(x) = \int P(x | \theta)P(\theta)d\theta$$

Bayesian Hierarchical model

Prior distribution of gene variances

$$\frac{1}{\sigma_g^2} | d_0, s_0^2 \sim \frac{1}{d_0 s_0^2} \chi_{d_0}^2$$

Hyperparameters (d_0, s_0^2)

Sampling distribution of sample variance for gene g

$$s_g^2 | \sigma_g^2 \sim \frac{\sigma_g^2}{d_g} \chi_{d_g}^2$$

σ_g^2 is a random variable

$$x_g^1, x_g^2, \dots, x_g^n \stackrel{ind}{\sim} N(\mu_g, \sigma_g^2)$$

expression is normally distributed
- hierarchy placed on variances

Estimate hyperparameters (d_0, s_0^2)

$$p(s^2 | d_0, s_0^2) = \int p(s^2 | \sigma^{-2}) p(\sigma^{-2}) d(\sigma^{-2}) \quad \text{Marginal distribution}$$



$$s^2 \sim s_0^2 F_{d_g, d_0}$$

Sample variances follow a scaled F-distribution

Estimate s_0 and d_0 ? `limma::fitFDist(x, df1)`

$$z_g = \log s_g^2$$

- z_g follows a Fisher's z-distribution
- hyperparameters are estimated by matching the theoretical mean and variance of the z-distribution to the observed sample mean and variance of z_g

The distributional result assumes that the typical prior gene-wise error variance s_0^2 and its associated degrees of freedom d_0 are known, which of course they are not. In practice, they will be estimated from the data itself (which is what the term empirical Bayes refers to).

These are examples of hyperparameters. In a full blown Bayesian approach, the user would specify *a priori*.

Moderate genewise variances

$$\sigma_g^2 | s_g^2 \sim \frac{d_0 s_0^2 + d_g s_g^2}{\chi_{d_0 + d_g}^2} \quad \text{Posterior distribution}$$

$$\tilde{s}_g^2 = s_{g[\text{moderated}]}^2 = \frac{d_g s_g^2 + d_0 s_0^2}{d_g + d_0} = \lambda s_g^2 + (1 - \lambda) s_0^2$$

$$\text{with } \lambda = \frac{d_g}{d_g + d_0} \in (0, 1)$$

`?limma::squeezeVar(var, df1)`

The posterior mean for gene-specific variance:

$$\tilde{s}_g^2 = \frac{d_0 s_0^2 + d_g s_g^2}{d_0 + d_g}$$

How to think about it:

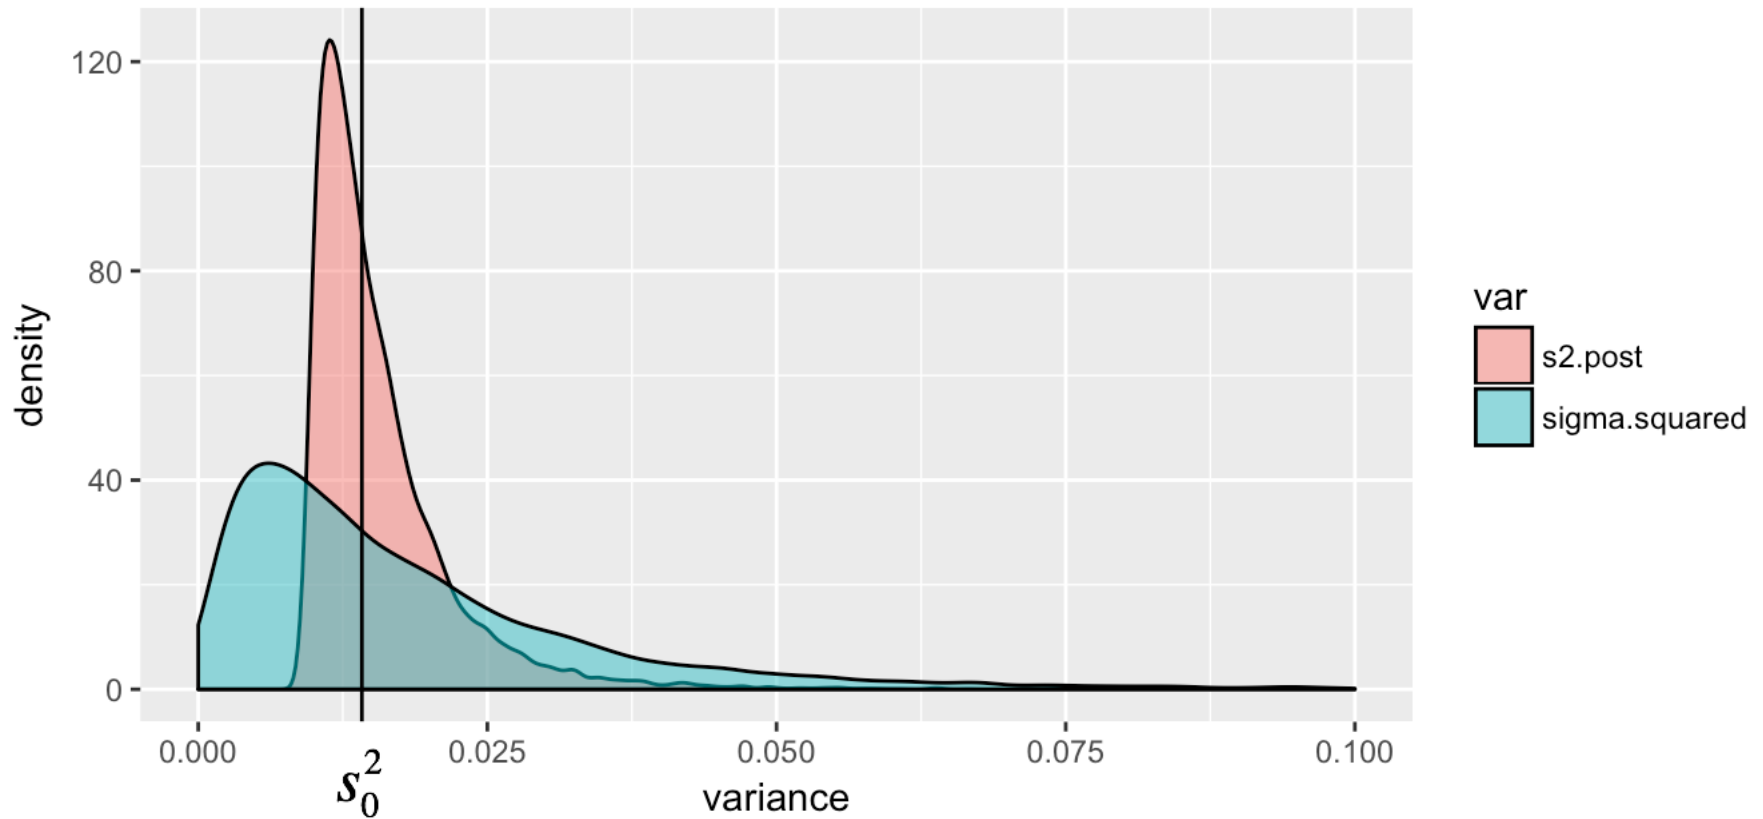
a weighted mean of the prior (indirect evidence) and the observed (direct evidence) gene-specific variances

$$\tilde{s}_g^2 = \frac{d_0}{d_0 + d_g} s_0^2 + \frac{d_g}{d_0 + d_g} s_g^2$$

More how to think about it:

“shrinking” the observed gene-specific variance towards the “typical” variance implied by the prior

Shrink variances – spikein95 data



`s2.prior` estimated prior value for σ^2

`s2.post` numeric vector giving the posterior values for σ^2

use this posterior mean to get a *moderated* t-statistic

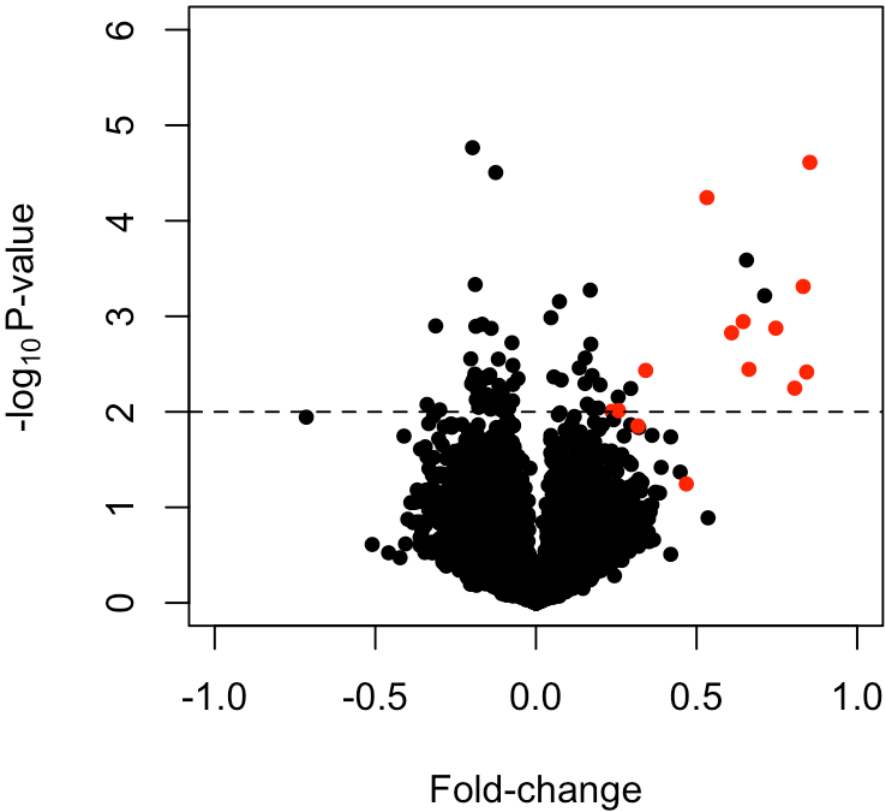
$$\tilde{t}_{gj} = \frac{\hat{\beta}_{gj}}{\tilde{s}_g \sqrt{v_j}}$$

under limma assumptions, we have the null distribution for this moderated t-statistic

$$\tilde{t}_{gj} \sim t_{d_0+d_g} \text{ under } H_0$$

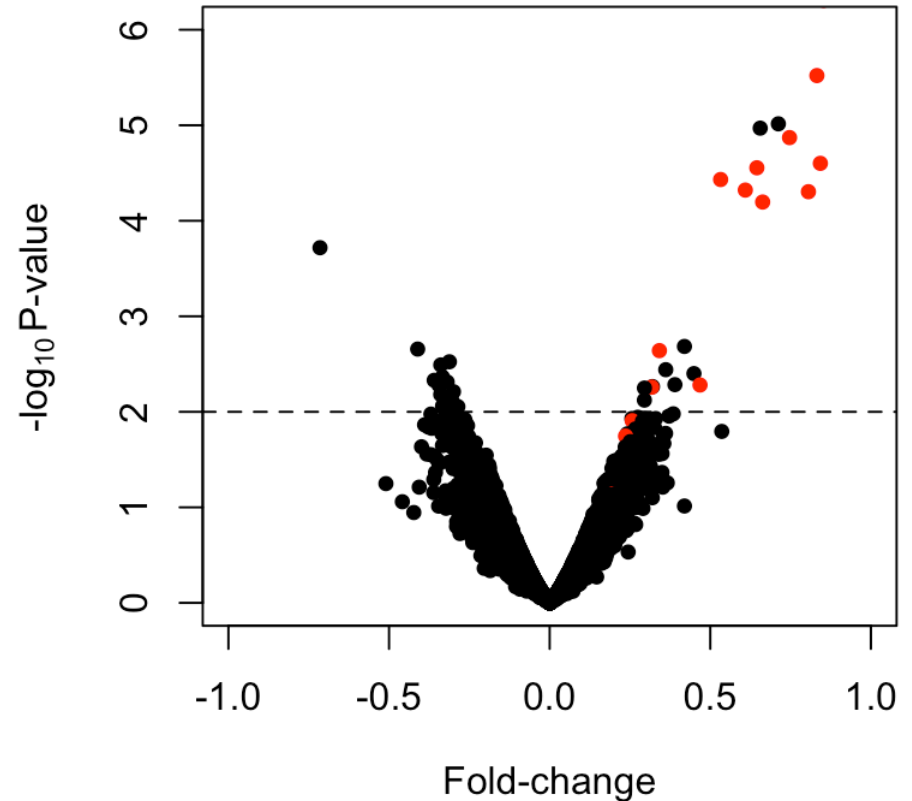
Revisit - spikein95

A) Volcano plot for t-test



1470 differentially expressed genes!! –
majority have very small variances

B) Volcano plot for moderated t-test



267 differentially expressed genes

side-by-side comparison of key quantities and results

“plain vanilla”

limma

estimated gene-wise residual variance

$$s_g^2 = \frac{1}{n-p} \hat{\boldsymbol{\varepsilon}}_g^T \hat{\boldsymbol{\varepsilon}}_g = \frac{1}{n-p} \sum_{i=1}^n \varepsilon_{gi}^2 \quad \tilde{s}_g^2 = \frac{d_0 s_0^2 + d_g s_g^2}{d_0 + d_g}$$

t-statistic for $H_0 : \beta_{gj} = 0$

$$t_{gj} = \frac{\hat{\beta}_{gj}}{s_g \sqrt{v_j}}$$

$$\tilde{t}_{gj} = \frac{\hat{\beta}_{gj}}{\tilde{s}_g \sqrt{v_j}}$$

distribution of the t-statistic when $\beta_{gj} = 0$

$$t_{gj} \sim t_{d_g}$$

$$\tilde{t}_{gj} \sim t_{d_0 + d_g}$$

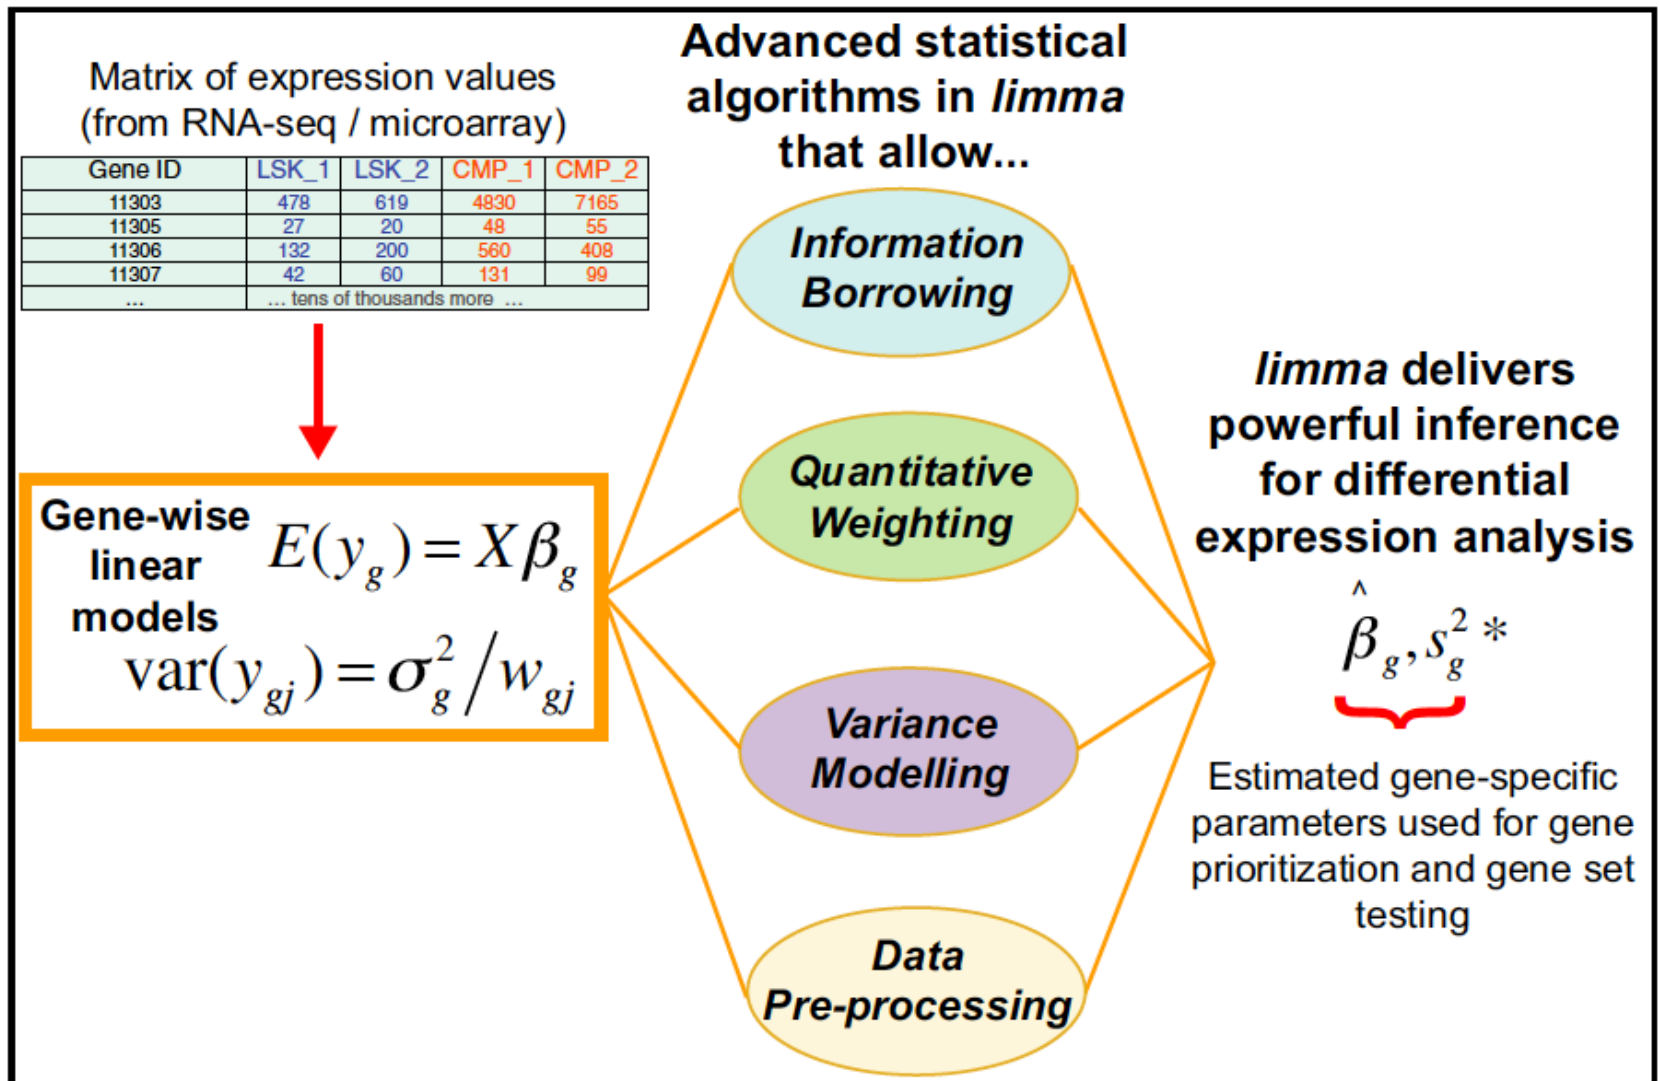
moderated variances will be “shrunk” towards the typical gene-wise variance, relative to raw sample residual variances

should result in fewer small variances and large t-stats

degrees of freedom for null distribution goes up, relative to default $d = n - p \rightarrow$ makes it closer to a standard normal \rightarrow makes tail probabilities smaller \rightarrow makes p-values smaller

overall, when all is well, limma will deliver statistical results that are more stable, more powerful

Linear Models for Microarrays and RNA-Seq Data (limma)



Smyth, Gordon K. (2004) “Linear Models and Empirical Bayes Methods for Assessing Differential Expression in Microarray Experiments,” Statistical Applications in Genetics and Molecular Biology: Vol. 3 : Iss. 1, Article 3.

DOI: 10.2202/1544-6115.1027

Available at: <http://www.bepress.com/sagmb/vol3/iss1/art3>

link no longer works now that SAGMB has been assimilated into the ~~Berg~~ deGruyter

But you should probably regard this as more definitive:

(Smyth describes as a “Reprint PDF with corrections”) and it’s dated 30 June 2009)

<http://www.statsci.org/smyth/pubs/ebayes.pdf>

<http://bioinf.wehi.edu.au/limma/>

<http://www.statsci.org/smyth/index.html>

Bioinformatics and Computational Biology Solutions Using R and Bioconductor -- [eBook](#) | Robert Gentleman, Vincent J. Carey, Wolfgang Huber, Rafael A. Irizarry, and Sandrine Dudoit, Springer 2005. Chapters 23 (limma: Linear Models for Microarray, by Smyth) and 14 (Analysis of Differential Gene Expression Studies, by Scholtens and von Heydebreck) are especially relevant.

limma:
Linear Models for Microarray and RNA-Seq Data
User's Guide

Gordon K. Smyth, Matthew Ritchie, Natalie Thorne,
James Wettenhall, Wei Shi and Yifang Hu
Bioinformatics Division, The Walter and Eliza Hall Institute
of Medical Research, Melbourne, Australia

First edition 2 December 2002

Last revised 16 October 2016

<http://www.bioconductor.org/packages/devel/bioc/vignettes/limma/inst/doc/usersguide.pdf>

specific sections I **strongly encourage** you to read

Ch. 8 Linear Models Overview

8.1 Introduction p. 35

8.2 Single-Channel Designs p. 36

Ch. 9 Single-Channel Experimental Designs

9.1 Introduction p. 40

9.2 Two Groups p.40

9.3 Several Groups p.42

9.4 Additive Models and Blocking p. 43

9.5 Interaction Models: 2 x 2 Factorial Designs p. 43*

Ch. 13 Statistics for Differential Expression

13.1 Summary Top-Tables p. 60

13.2 Fitted Model Objects p. 61

13.3 Multiple Testing Across Contrasts p. 62

Ch. 16 and Ch. 17: Case Studies

15 RNA-seq Data	69
15.1 Introduction	69
15.2 Making a count matrix	69
15.3 Normalization and filtering	69
15.4 Differential expression: limma-trend	70
15.5 Differential expression: voom	70
15.6 Voom with sample quality weights	71
15.7 Differential splicing	73






Chapter 7

Filtering

Note that filtering methods involving variances should not be used. The limma algorithm analyses the spread of the genewise variances. Any filtering method based on genewise variances will change the distribution of variances, will interfere with the limma algorithm and hence will give poor results.

Functions that make your life easier:

functions

model.matrix		Takes in your “factors” and makes a design matrix
lmFit		Fits the linear model to all genes (each gene separately) – replace gene with “feature” depending on your data.
makeContrasts		Create the contrast matrix C that you desire
eBayes		Use output of linear regression to compute moderated t statistics
topTable		Query your results; sort your pvalues; sort genes; Adjust for multiple comparisons etc

23.8 Several groups

The above approaches for two groups extend easily to any number of groups. Suppose that three RNA targets are to be compared using Affymetrix arrays. Suppose that the three targets are called “RNA1,” “RNA2,” and “RNA3” and that the column `targets$Target` indicates which one was hybridized to each array. An appropriate design matrix can be created using

```
> f <- factor(targets$Target, levels = c("RNA1",  
+   "RNA2", "RNA3"))
```

x is categorical, a factor
specifies 3 groups

make design matrix, request “cell means” parametrization

```
> design <- model.matrix(~0 + f)  
> colnames(design) <- c("RNA1", "RNA2", "RNA3")
```

To make all pair-wise comparisons between the three groups, one could proceed

fit linear model using least squares

```
> fit <- lmFit(eset, design)  
> contrast.matrix <- makeContrasts(RNA2 - RNA1,  
+   RNA3 - RNA2, RNA3 - RNA1, levels = design)  
> fit2 <- contrasts.fit(fit, contrast.matrix)  
> fit2 <- eBayes(fit2)
```

specify contrasts of interest;
here, all three pairwise comparisons

moderate the test stats

A list of top genes for RNA2 versus RNA1 can be obtained from

```
> topTable(fit2, coef = 1, adjust = "fdr")
```

produce FDR-adjusted p-values, a la Benjamini-Hochberg, for gene-wise test of H_0 : contrast #1 = 0, sort in order of statistical significance, and report the top hits

```
> system.time(jFit <- lmFit(prDat, jDesMat))  
   user  system elapsed  
0.345   0.113   0.459
```

using limma's `lmFit()` function to perform two-way ANOVA for ~30K probesets

this takes a trivial amount of time

the hard parts for the analyst are choosing the model, choosing how to parametrize it and digesting the results

novices will be surprised what a non-issue the number of genes, number of samples is

wise words (I find) relevant to science, statistical analysis, frequentism vs. Bayesianism, pragmatic approaches vs. mathematically pure ones,

All models are ~~wrong~~, but some are useful. (George E. P. Box)
simplification

"An approximate answer to the right problem is worth a good deal more than an exact answer to an approximate problem." John Tukey

"Absolute certainty is a privilege of uneducated minds and fanatics. It is, for scientific folk, an unattainable ideal." Cassius J. Keyser



Jenny Bryan
@JennyBryan

Following

All models are wrong, so why not start with one you actually understand?

RETWEETS
168

LIKES
319



2:09 PM - 1 Oct 2016



11



168



319

```

> jDesMat <- model.matrix(~ gType * devStage, prDes)
> ## ridiculous machination to print a version to screen with small
> ## variable names
> foo <- jDesMat
> colnames(foo) <- paste0("X", formatC(seq_len(ncol(jDesMat)), width = 2, flag = "0"))
> cbind(prDes, foo)

```

	sample	devStage	gType	X01	X02	X03	X04	X05	X06	X07	X08	X09	X10
Sample_20	20	E16	wt	1	0	0	0	0	0	0	0	0	0
Sample_21	21	E16	wt	1	0	0	0	0	0	0	0	0	0
Sample_22	22	E16	wt	1	0	0	0	0	0	0	0	0	0
Sample_23	23	E16	wt	1	0	0	0	0	0	0	0	0	0
Sample_16	16	E16	NrlKO	1	1	0	0	0	0	0	0	0	0
Sample_17	17	E16	NrlKO	1	1	0	0	0	0	0	0	0	0
Sample_6	6	E16	NrlKO	1	1	0	0	0	0	0	0	0	0
Sample_24	24	P2	wt	1	0	1	0	0	0	0	0	0	0
Sample_25	25	P2	wt	1	0	1	0	0	0	0	0	0	0
Sample_26	26	P2	wt	1	0	1	0	0	0	0	0	0	0
Sample_27	27	P2	wt	1	0	1	0	0	0	0	0	0	0
Sample_14	14	P2	NrlKO	1	1	1	0	0	0	1	0	0	0
Sample_3	3	P2	NrlKO	1	1	1	0	0	0	1	0	0	0
Sample_5	5	P2	NrlKO	1	1	1	0	0	0	1	0	0	0
Sample_8	8	P2	NrlKO	1	1	1	0	0	0	1	0	0	0
Sample_28	28	P6	wt	1	0	0	1	0	0	0	0	0	0
Sample_29	29	P6	wt	1	0	0	1	0	0	0	0	0	0
Sample_30	30	P6	wt	1	0	0	1	0	0	0	0	0	0
Sample_31	31	P6	wt	1	0	0	1	0	0	0	0	0	0
Sample_1	1	P6	NrlKO	1	1	0	1	0	0	0	1	0	0
Sample_10	10	P6	NrlKO	1	1	0	1	0	0	0	1	0	0
Sample_4	4	P6	NrlKO	1	1	0	1	0	0	0	1	0	0
Sample_7	7	P6	NrlKO	1	1	0	1	0	0	0	1	0	0
Sample_32	32	P10	wt	1	0	0	0	1	0	0	0	0	0
Sample_33	33	P10	wt	1	0	0	0	1	0	0	0	0	0
Sample_34	34	P10	wt	1	0	0	0	1	0	0	0	0	0
Sample_35	35	P10	wt	1	0	0	0	1	0	0	0	0	0
Sample_13	13	P10	NrlKO	1	1	0	0	1	0	0	0	1	0
Sample_15	15	P10	NrlKO	1	1	0	0	1	0	0	0	1	0
Sample_18	18	P10	NrlKO	1	1	0	0	1	0	0	0	1	0
Sample_19	19	P10	NrlKO	1	1	0	0	1	0	0	0	1	0
Sample_36	36	4_weeks	wt	1	0	0	0	0	1	0	0	0	0
Sample_37	37	4_weeks	wt	1	0	0	0	0	1	0	0	0	0
Sample_38	38	4_weeks	wt	1	0	0	0	0	1	0	0	0	0
Sample_39	39	4_weeks	wt	1	0	0	0	0	1	0	0	0	0
Sample_11	11	4_weeks	NrlKO	1	1	0	0	0	1	0	0	0	1
Sample_12	12	4_weeks	NrlKO	1	1	0	0	0	1	0	0	0	1
Sample_2	2	4_weeks	NrlKO	1	1	0	0	0	1	0	0	0	1
Sample_9	9	4_weeks	NrlKO	1	1	0	0	0	1	0	0	0	1

human- and
computer-
readable info on
factor levels for
each sample =
prDes

lm.fit- and lmFit-
ready encoding
of factor levels
for each sample
=
a design matrix =
X or X_g

θ τ_{NrlKO} β_{P2} , etc $(\tau\beta)_{NrlKO,P2}$, etc

```

> cbind(colnames(jDesMat))
      [,1]
[1,] "(Intercept)"
[2,] "gTypeNrlKO"
[3,] "devStageP2"
[4,] "devStageP6"
[5,] "devStageP10"
[6,] "devStage4_weeks"
[7,] "gTypeNrlKO:devStageP2"
[8,] "gTypeNrlKO:devStageP6"
[9,] "gTypeNrlKO:devStageP10"
[10,] "gTypeNrlKO:devStage4_weeks"

```

```
> jDesMat <- model.matrix(~ gType * devStage, prDes)
> jFit <- lmFit(prDat, jDesMat)

> ebFit <- eBayes(jFit)
```

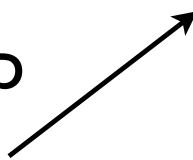
```
> summary(jFit)
```

	Length	Class	Mode
coefficients	299490	-none-	numeric
rank	1	-none-	numeric
assign	10	-none-	numeric
qr	5	qr	list
df.residual	29949	-none-	numeric
sigma	29949	-none-	numeric
cov.coefficients	100	-none-	numeric
stdev.unscaled	299490	-none-	numeric
pivot	10	-none-	numeric
genes	1	data.frame	list
Amean	29949	-none-	numeric
method	1	-none-	character
design	390	-none-	numeric

```
> summary(ebFit)
```

	Length	Class	Mode
coefficients	299490	-none-	numeric
rank	1	-none-	numeric
assign	10	-none-	numeric
qr	5	qr	list
df.residual	29949	-none-	numeric
sigma	29949	-none-	numeric
cov.coefficients	100	-none-	numeric
stdev.unscaled	299490	-none-	numeric
pivot	10	-none-	numeric
genes	1	data.frame	list
Amean	29949	-none-	numeric
method	1	-none-	character
design	390	-none-	numeric
df.prior	1	-none-	numeric
s2.prior	1	-none-	numeric
var.prior	10	-none-	numeric
proportion	1	-none-	numeric
s2.post	29949	-none-	numeric
t	299490	-none-	numeric
df.total	29949	-none-	numeric
p.value	299490	-none-	numeric
lods	299490	-none-	numeric
F	29949	-none-	numeric
F.p.value	29949	-none-	numeric

see all this new stuff?
topTable() will help
you use it to find
interesting genes



limma is designed to help you out
AFTER you've applied eBayes()

limma workflow

fit a separate linear model for
each response, e.g. gene

`lmFit(...)`

fitted models

apply an Empirical Bayes
procedure for moderating
estimates of error variance

`eBayes(...)`

extract estimated parameters
or p-values or ...
compare big models to small
etc etc

`topTable(...)`


```
> jDesMat <- model.matrix(~ gType * devStage, prDes)
> jFit <- lmFit(prDat, jDesMat)
> ebFit <- eBayes(jFit)
```

```
> topTable(ebFit)

      ID X.Intercept.  gTypeNrlKO devStageP2 devStageP6
1438940_x_at 1438940_x_at    12.8625  0.05750000    0.0850   0.1325
1436884_x_at 1436884_x_at    12.9275  0.05916667    0.1775   0.3225
1456736_x_at 1456736_x_at    12.3225 -0.07583333    0.1625   0.3050
1455897_x_at 1455897_x_at    13.0575  0.01916667   -0.0150   0.1100
1451240_a_at 1451240_a_at    12.9975 -0.03083333    0.3100   0.2800
1454613_at   1454613_at     12.4675 -0.28750000   -0.1075  -0.0500
1450084_s_at 1450084_s_at    12.6350 -0.04500000    0.0825   0.0525
1437192_x_at 1437192_x_at    12.9425  0.07750000    0.2750   0.3000
1449676_at   1449676_at     12.7075 -0.05750000   -0.0750  -0.1525
1438657_x_at 1438657_x_at    12.7825  0.15083333    0.1400   0.1250

      devStageP10 devStage4_weeks  gTypeNrlKO.devStageP2
1438940_x_at      0.3425          0.3500          0.01750000
1436884_x_at      0.0300          0.0250         -0.24166667
1456736_x_at      0.2075          0.0725          0.03583333
1455897_x_at      0.4325          0.4775          0.09083333
1451240_a_at      0.2800         -0.3700          0.23083333
1454613_at     -0.1025         -0.3825          0.15500000
1450084_s_at      0.1725          0.2600          0.13000000
1437192_x_at      0.2925         -0.0050         -0.19750000
1449676_at     -0.1725         -0.5075          0.17250000
1438657_x_at     -0.1850         -0.4500         -0.30083333

      gTypeNrlKO.devStageP6  gTypeNrlKO.devStageP10
1438940_x_at      0.05500000         -0.12000000
1436884_x_at     -0.37666667         -0.10166667
1456736_x_at     -0.02416667         -0.14416667
1455897_x_at      0.19583333         -0.06666667
1451240_a_at      0.28833333          0.18583333
1454613_at      0.15000000          0.27250000
1450084_s_at      0.05000000          0.06250000
1437192_x_at     -0.23750000         -0.21500000
1449676_at      0.28500000          0.28250000
1438657_x_at     -0.29833333          0.04166667

      gTypeNrlKO.devStage4_weeks  AveExpr      F      P.Value
1438940_x_at    -0.132500000    13.05872  63728.49  5.063198e-66
1436884_x_at    -0.009166667   12.99538  52673.21  1.077659e-64
1456736_x_at      0.060833333   12.43154  51040.87  1.786135e-64
1455897_x_at    -0.094166667   13.28590  49456.68  2.962710e-64
1451240_a_at      0.720833333   13.23128  48588.27  3.937053e-64
1454613_at      0.430000000   12.29897  43799.55  2.081658e-63
1450084_s_at    -0.010000000   12.75333  43532.81  2.296095e-63
1437192_x_at      0.030000000   13.09359  42553.50  3.308144e-63
1449676_at      0.515000000   12.62205  42311.46  3.625302e-63
1438657_x_at      0.086666667   12.73179  42145.13  3.861878e-63

      adj.P.Val
1438940_x_at  1.516377e-61
1436884_x_at  1.613741e-60
1456736_x_at  1.783099e-60
1455897_x_at  2.218255e-60
1451240_a_at  2.358216e-60
1454613_at   9.823679e-60
1450084_s_at  9.823679e-60
1437192_x_at  1.156594e-59
1449676_at   1.156594e-59
1438657_x_at  1.156594e-59
```

however, you can't just use
topTable() on auto-pilot

you still must know and
describe what you consider
a “hit” to be

```
topTable(fit, coef=NULL, number=10, genelist=fit$genes, adjust.method="BH",  
         sort.by="B", resort.by=NULL, p.value=1, lfc=0, confint=FALSE)
```

```
topTableF(fit, number=10, genelist=fit$genes, adjust.method="BH",  
          sort.by="F", p.value=1, lfc=0)
```

`coef`: column number or column name specifying which coefficient or contrast of the linear model is of interest. For 'topTable', can also be a vector of column subscripts, in which case the gene ranking is by F-statistic for that set of contrasts.

'topTableF' ranks genes on the basis of moderated F-statistics for one or more coefficients. If 'topTable' is called with 'coef' has length greater than 1, then the specified columns will be extracted from 'fit' and 'topTableF' called on the result. 'topTable' with 'coef=NULL' is the same as 'topTableF', unless the fitted model 'fit' has only one column.

coef argument is where you specify what it is you want to test for equality with zero

Recent limma feature

Estimating gene-specific variance priors!

The Annals of Applied Statistics

2016, Vol. 10, No. 2, 946–963

DOI: 10.1214/16-AOAS920

© Institute of Mathematical Statistics, 2016

ROBUST HYPERPARAMETER ESTIMATION PROTECTS AGAINST HYPERVARIABLE GENES AND IMPROVES POWER TO DETECT DIFFERENTIAL EXPRESSION¹

BY BELINDA PHIPSON^{*}, STANLEY LEE^{†,‡}, IAN J. MAJEWSKI^{†,‡},
WARREN S. ALEXANDER^{†,‡} AND GORDON K. SMYTH^{†,‡}

`eBayes(fit, robust=TRUE)`